

#PROJECT- Classification model for ABG motors using predictive analytics to enter the Indian Market
#Aayusha Chakraborty

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report, roc_auc_score
```

First we will load the excel file
file_path = r"C:\Users\chakr\OneDrive\Documents\
Capstone.datasciencel.xlsx"

Loading specific sheets into DataFrames
indian_df = pd.read_excel(file_path, sheet_name='indiandataset ')
japanese_df = pd.read_excel(file_path, sheet_name='japanesedataset ')

Handle missing values (if any)
indian_df = indian_df.dropna() *# Drop rows with missing values in the Indian dataset*
japanese_df = japanese_df.dropna() *# Drop rows with missing values in the Japanese dataset*

Display the first few rows of each DataFrame to verify
print("Indian Dataset after dropping missing values:")
print(indian_df.head())

print("\nJapanese Dataset after dropping missing values:")
print(japanese_df.head())

Indian Dataset after dropping missing values:

	ID	CURR_AGE	GENDER	ANN_INCOME	DT_MAINT
0	20710B05XL	54	M	1425390	4/20/2018
1	89602T51HX	47	M	1678954	2018-08-06 00:00:00
2	70190Z52IP	60	M	931624	7/31/2017
3	25623V15MU	55	F	1106320	7/31/2017
4	36230I68CE	32	F	748465	1/27/2019

Japanese Dataset after dropping missing values:

Empty DataFrame

Columns: [ID, CURR_AGE, GENDER, ANN_INCOME, AGE_CAR, PURCHASE,
Unnamed: 6, Unnamed: 7, Unnamed: 8]

Index: []

```
indian_df = pd.read_excel(file_path, sheet_name='indiandataset ')  
japanese_df = pd.read_excel(file_path, sheet_name='japanesedataset ')
```

```

# Strip any leading or trailing spaces from column names
indian_df.columns = indian_df.columns.str.strip()
japanese_df.columns = japanese_df.columns.str.strip()

# Verify column names
print("Indian Dataset Columns:", indian_df.columns)
print("Japanese Dataset Columns:", japanese_df.columns)

# We have done the segmentation of the age_car provided in the dataset
into four categories
def segment_age_car(age_car):
    if age_car < 200:
        return 1
    elif 200 <= age_car <= 360:
        return 2
    elif 360 < age_car <= 500:
        return 3
    else:
        return 4

# Applied the segmentation to the 'AGE_CAR' column in the Japanese
dataset
if 'AGE_CAR' in japanese_df.columns:
    japanese_df['age_car_segment'] =
japanese_df['AGE_CAR'].apply(segment_age_car)
    print(japanese_df[['AGE_CAR', 'age_car_segment']].head())
else:
    print("Column 'AGE_CAR' does not exist in the Japanese dataset.")

Indian Dataset Columns: Index(['ID', 'CURR_AGE', 'GENDER',
'ANN_INCOME', 'DT_MAINT'], dtype='object')
Japanese Dataset Columns: Index(['ID', 'CURR_AGE', 'GENDER',
'ANN_INCOME', 'AGE_CAR', 'PURCHASE',
'Unnamed: 6', 'Unnamed: 7', 'Unnamed: 8'],
dtype='object')
  AGE_CAR  age_car_segment
0      439                3
1      283                2
2      390                3
3      475                3
4      497                3

# We have the annual income of the customers in currency yen, so we
have converted into inr with a conversion rate of 0.64 as of year 2019
conversion_rate = 0.64

# Convert 'ANN_INCOME' from Yen to INR
japanese_df['ANN_INCOME_INR'] = japanese_df['ANN_INCOME'] *
conversion_rate

```

```
# Display the first few rows to check the conversion
print(japanese_df[['ANN_INCOME', 'ANN_INCOME_INR']].head())
```

	ANN_INCOME	ANN_INCOME_INR
0	445344	285020.16
1	107634	68885.76
2	502787	321783.68
3	585664	374824.96
4	705723	451662.72

```
# Calculated maximum and minimum values to calculate the range
```

```
max_income_inr = japanese_df['ANN_INCOME_INR'].max()
```

```
min_income_inr = japanese_df['ANN_INCOME_INR'].min()
```

```
# Print the results
```

```
print(f"Maximum Annual Income in INR: {max_income_inr}")
```

```
print(f"Minimum Annual Income in INR: {min_income_inr}")
```

Maximum Annual Income in INR: 511981.44

Minimum Annual Income in INR: 44856.96

```
# classify income groups based on the INR values for japanese dataset
```

```
def classify_income(income_inr):
    if income_inr <= 100000:
        return 1 # Low income
    elif 100000 < income_inr <= 200000:
        return 2 # Lower-middle income
    elif 200000 < income_inr <= 300000:
        return 3 # Upper-middle income
    else:
        return 4 # High income
```

```
# Apply the classification to the Japanese dataset
```

```
japanese_df['income_group'] =
```

```
japanese_df['ANN_INCOME_INR'].apply(classify_income)
```

```
# Check the classification
```

```
print(japanese_df[['ANN_INCOME_INR', 'income_group']].head())
```

	ANN_INCOME_INR	income_group
0	285020.16	3
1	68885.76	1
2	321783.68	4
3	374824.96	4
4	451662.72	4

```
# Calculate maximum and minimum age in the Japanese dataset
```

```
max_age_japanese = japanese_df['CURR_AGE'].max()
```

```
min_age_japanese = japanese_df['CURR_AGE'].min()
```

```
# Print the results
```

```
print(f"Maximum Age in Japanese Dataset: {max_age_japanese}")
print(f"Minimum Age in Japanese Dataset: {min_age_japanese}")
```

```
Maximum Age in Japanese Dataset: 65
Minimum Age in Japanese Dataset: 25
```

```
# Calculate maximum and minimum age in the indian dataset
max_age_indian = indian_df['CURR_AGE'].max()
min_age_indian = indian_df['CURR_AGE'].min()
```

```
# Print the results
print(f"Maximum Age in indian Dataset: {max_age_indian}")
print(f"Minimum Age in indian Dataset: {min_age_indian}")
```

```
Maximum Age in indian Dataset: 65
Minimum Age in indian Dataset: 25
```

```
# Here we have classified age into four categories:
def classify_age_group(age):
    if 25 <= age < 35:
        return 1
    elif 35 <= age < 45:
        return 2
    elif 45 <= age < 55:
        return 3
    elif 55 <= age <= 65:
        return 4
    else:
        return None # This handles any unexpected values
```

```
# Apply the classification to the 'CURR_AGE' column
indian_df['age_group'] =
indian_df['CURR_AGE'].apply(classify_age_group)
```

```
# Check the new classification
print(indian_df[['CURR_AGE', 'age_group']].head())
```

	CURR_AGE	age_group
0	54	3
1	47	3
2	60	4
3	55	4
4	32	1

```
def classify_age_group(age):
    if 25 <= age < 35:
        return 1
    elif 35 <= age < 45:
        return 2
    elif 45 <= age < 55:
        return 3
```

```

    elif 55 <= age <= 65:
        return 4
    else:
        return None # This handles any unexpected values

# Apply the classification to the 'CURR_AGE' column
japanese_df['age_group'] =
japanese_df['CURR_AGE'].apply(classify_age_group)

# Check the new classification
print(japanese_df[['CURR_AGE', 'age_group']].head())

```

	CURR_AGE	age_group
0	50	3
1	35	2
2	59	4
3	43	2
4	39	2

```

# Here we have converted the string value into binary of the gender
column
def convert_gender_indian(gender):
    if gender == 'M':
        return 0
    elif gender == 'F':
        return 1
    else:
        return None # Handle unexpected values

# Apply the conversion to the Indian dataset
indian_df['gender_numeric'] =
indian_df['GENDER'].apply(convert_gender_indian)

# Define a function to convert gender into numerical values for the
Japanese dataset
def convert_gender_japanese(gender):
    if gender == 'M':
        return 0
    elif gender == 'F':
        return 1
    else:
        return None # Handle unexpected values

# Apply the conversion to the Japanese dataset
japanese_df['gender_numeric'] =
japanese_df['GENDER'].apply(convert_gender_japanese)

# Check the new columns
print("Indian Dataset - GENDER and gender_numeric:")
print(indian_df[['GENDER', 'gender_numeric']].head())

```

```
print("Japanese Dataset - GENDER and gender_numeric:")
print(japanese_df[['GENDER', 'gender_numeric']].head())
```

Indian Dataset - GENDER and gender_numeric:

	GENDER	gender_numeric
0	M	0
1	M	0
2	M	0
3	F	1
4	F	1

Japanese Dataset - GENDER and gender_numeric:

	GENDER	gender_numeric
0	M	0
1	M	0
2	F	1
3	M	0
4	F	1

Calculate the maximum and minimum of the annual income

```
max_income = indian_df['ANN_INCOME'].max()
min_income = indian_df['ANN_INCOME'].min()
```

```
print(f"Maximum Annual Income: {max_income}")
print(f"Minimum Annual Income: {min_income}")
```

```
Maximum Annual Income: 1999989
Minimum Annual Income: 300033
```

Convert DT_MAINT to datetime format (if not already in datetime)

```
indian_df['DT_MAINT'] = pd.to_datetime(indian_df['DT_MAINT'],
errors='coerce')
```

Define the reference date (7th July 2019)

```
reference_date = pd.to_datetime('2019-07-07')
```

Calculate the age of the car by subtracting DT_MAINT from the reference date

The result will be in days, so we can divide by 365 to get years

```
indian_df['age_car'] = (reference_date -
indian_df['DT_MAINT']).dt.days
```

Display the updated dataframe with the new age_car column

```
print(indian_df[['ID', 'DT_MAINT', 'age_car']].head())
```

	ID	DT_MAINT	age_car
0	20710B05XL	2018-04-20	443
1	89602T51HX	2018-08-06	335
2	70190Z52IP	2017-07-31	706
3	25623V15MU	2017-07-31	706
4	36230I68CE	2019-01-27	161

```

# Define a function to segment the age_car values
def segment_age_car(age):
    if age < 200:
        return 1
    elif 200 <= age <= 360:
        return 2
    elif 360 < age <= 500:
        return 3
    else:
        return 4

# Apply the function to the age_car column and create a new column
with the segments
indian_df['age_car_segment'] =
indian_df['age_car'].apply(segment_age_car)

# Display the updated dataframe with the age_car and age_car_segment
columns
print(indian_df[['ID', 'DT_MAINT', 'age_car',
'age_car_segment']].head())

```

	ID	DT_MAINT	age_car	age_car_segment
0	20710B05XL	2018-04-20	443	3
1	89602T51HX	2018-08-06	335	2
2	70190Z52IP	2017-07-31	706	4
3	25623V15MU	2017-07-31	706	4
4	36230I68CE	2019-01-27	161	1

```

# Define a function to classify current age into age groups
def classify_age_group(age):
    if 25 <= age < 35:
        return 1
    elif 35 <= age < 45:
        return 2
    elif 45 <= age < 55:
        return 3
    elif 55 <= age <= 65:
        return 4

# Apply the function to the CURR_AGE column and create a new column
'age_group'
indian_df['age_group'] =
indian_df['CURR_AGE'].apply(classify_age_group)

# Display the updated dataframe with CURR_AGE and age_group columns
print(indian_df[['ID', 'CURR_AGE', 'age_group']].head())

```

	ID	CURR_AGE	age_group
0	20710B05XL	54	3

1	89602T51HX	47	3
2	70190Z52IP	60	4
3	25623V15MU	55	4
4	36230I68CE	32	1

```
print(indian_df.head(10))
```

	ID	CURR_AGE	GENDER	ANN_INCOME	DT_MAINT	age_group \
0	20710B05XL	54	M	1425390	2018-04-20	3
1	89602T51HX	47	M	1678954	2018-08-06	3
2	70190Z52IP	60	M	931624	2017-07-31	4
3	25623V15MU	55	F	1106320	2017-07-31	4
4	36230I68CE	32	F	748465	2019-01-27	1
5	11264G01HZ	48	F	1051927	2018-11-24	3
6	74250S23U0	26	F	1076402	2018-09-22	1
7	26735J66DB	45	F	1481949	2018-05-04	3
8	93404P60ED	55	M	1725607	2018-02-01	4
9	56557A36QV	64	F	312323	2018-04-23	4

	gender_numeric	age_car	age_car_segment
0	0	443	3
1	0	335	2
2	0	706	4
3	1	706	4
4	1	161	1
5	1	225	2
6	1	288	2
7	1	429	3
8	0	521	4
9	1	440	3

```
# Assuming `indian_df` is your DataFrame with the 'ANN_INCOME' column
```

```
# Step 1: Calculate mean and standard deviation of the 'ANN_INCOME' column
```

```
mean_income = indian_df['ANN_INCOME'].mean()
std_income = indian_df['ANN_INCOME'].std()
```

```
# Step 2: Apply standardization formula (x - mean) / std
```

```
indian_df['income_standardized'] =
indian_df['ANN_INCOME'].apply(lambda x: (x - mean_income) /
std_income)
```

```
# Step 3: Check the result
```

```
print(indian_df[['ANN_INCOME', 'income_standardized']].head())
```

	ANN_INCOME	income_standardized
0	1425390	0.692730
1	1678954	1.327512
2	931624	-0.543383

3	1106320	-0.106042
4	748465	-1.001910

#similarly we will do for japanese dataset

Step 1: Calculate mean and standard deviation of the 'ANN_INCOME' column

```
mean_income = japanese_df['ANN_INCOME'].mean()
```

```
std_income = japanese_df['ANN_INCOME'].std()
```

Step 2: Apply standardization formula $(x - \text{mean}) / \text{std}$

```
japanese_df['income_standardized'] =
```

```
japanese_df['ANN_INCOME'].apply(lambda x: (x - mean_income) /  
std_income)
```

Step 3: Check the result

```
print(japanese_df[['ANN_INCOME', 'income_standardized']].head())
```

	ANN_INCOME	income_standardized
0	445344	0.490809
1	107634	-1.437759
2	502787	0.818849
3	585664	1.292137
4	705723	1.977760

Assume the 'DT_MAINT' column is already present in your dataset.

```
indian_df.columns = indian_df.columns.str.strip()
```

```
print("Indian Dataset Columns:", indian_df.columns)
```

Step 1: Convert 'DT_MAINT' to datetime format

```
indian_df['DT_MAINT'] = pd.to_datetime(indian_df['DT_MAINT'])
```

Step 2: Calculate the age of the car

```
reference_date = pd.to_datetime('2019-07-07') # The fixed date for  
calculation
```

```
indian_df['age_car'] = (reference_date -  
indian_df['DT_MAINT']).dt.days
```

Step 3: Function to categorize age_car

```
def categorize_age_car(age_car):
```

```
    if age_car < 200:
```

```
        return 1
```

```
    elif 200 <= age_car <= 360:
```

```
        return 2
```

```
    elif 360 < age_car <= 500:
```

```
        return 3
```

```
    else:
```

```
        return 4
```

Apply the categorization function

```
indian_df['age_car_segment'] =
```

```
indian_df['age_car'].apply(categorize_age_car)
```

```
Indian Dataset Columns: Index(['ID', 'CURR_AGE', 'GENDER',
                              'ANN_INCOME', 'DT_MAINT', 'age_group',
                              'gender_numeric', 'age_car', 'age_car_segment',
                              'income_standardized'],
                              dtype='object')
```

```
print(indian_df.head(20))
```

	ID	CURR_AGE	GENDER	ANN_INCOME	DT_MAINT	age_group	\
0	20710B05XL	54	M	1425390	2018-04-20	3	
1	89602T51HX	47	M	1678954	2018-08-06	3	
2	70190Z52IP	60	M	931624	2017-07-31	4	
3	25623V15MU	55	F	1106320	2017-07-31	4	
4	36230I68CE	32	F	748465	2019-01-27	1	
5	11264G01HZ	48	F	1051927	2018-11-24	3	
6	74250S23U0	26	F	1076402	2018-09-22	1	
7	26735J66DB	45	F	1481949	2018-05-04	3	
8	93404P60ED	55	M	1725607	2018-02-01	4	
9	56557A36QV	64	F	312323	2018-04-23	4	
10	38353F50LZ	53	M	546574	2019-05-06	3	
11	54684T21RX	44	F	1203691	2017-12-07	2	
12	46929E04HS	59	F	724688	2019-06-22	4	
13	20647X82EQ	27	F	975130	2019-05-03	1	
14	34956P25RT	57	F	1422399	2019-04-14	4	
15	07090V20JQ	40	F	1558045	2018-06-02	2	
16	78392T89DQ	33	M	669737	2019-05-31	1	
17	07257K04CB	57	F	774593	2018-12-31	4	
18	65658K80PS	59	F	993201	2019-05-03	4	
19	69803K32CS	42	F	1050793	2016-10-21	2	

	gender_numeric	age_car	age_car_segment	income_standardized
0	0	443	3	0.692730
1	0	335	2	1.327512
2	0	706	4	-0.543383
3	1	706	4	-0.106042
4	1	161	1	-1.001910
5	1	225	2	-0.242212
6	1	288	2	-0.180940
7	1	429	3	0.834322
8	0	521	4	1.444305
9	1	440	3	-2.093765
10	0	62	1	-1.507332
11	1	577	4	0.137720
12	1	15	1	-1.061434
13	1	65	1	-0.434468
14	1	84	1	0.685243
15	1	400	3	1.024824
16	0	37	1	-1.199001
17	1	188	1	-0.936500

18	1	65	1	-0.389229
19	1	989	4	-0.245050

```
import pandas as pd
```

```
# Step 1: Standardize the ANN_INCOME column
```

```
mean_income = indian_df['ANN_INCOME'].mean()
```

```
std_income = indian_df['ANN_INCOME'].std()
```

```
indian_df['ANN_INCOME'] = (indian_df['ANN_INCOME'] - mean_income) /  
std_income
```

```
# Step 2: Replace DT_MAINT with age_car_segment
```

```
indian_df['DT_MAINT'] = indian_df['age_car_segment']
```

```
# Display the updated DataFrame with the first 50 rows
```

```
print(indian_df.head(50))
```

	ID	CURR_AGE	GENDER	ANN_INCOME	DT_MAINT	age_group	\
0	20710B05XL	54	M	0.692730	3	3	
1	89602T51HX	47	M	1.327512	2	3	
2	70190Z52IP	60	M	-0.543383	4	4	
3	25623V15MU	55	F	-0.106042	4	4	
4	36230I68CE	32	F	-1.001910	1	1	
5	11264G01HZ	48	F	-0.242212	2	3	
6	74250S23U0	26	F	-0.180940	2	1	
7	26735J66DB	45	F	0.834322	3	3	
8	93404P60ED	55	M	1.444305	4	4	
9	56557A36QV	64	F	-2.093765	3	4	
10	38353F50LZ	53	M	-1.507332	1	3	
11	54684T21RX	44	F	0.137720	4	2	
12	46929E04HS	59	F	-1.061434	1	4	
13	20647X82EQ	27	F	-0.434468	1	1	
14	34956P25RT	57	F	0.685243	1	4	
15	07090V20JQ	40	F	1.024824	3	2	
16	78392T89DQ	33	M	-1.199001	1	1	
17	07257K04CB	57	F	-0.936500	1	4	
18	65658K80PS	59	F	-0.389229	1	4	
19	69803K32CS	42	F	-0.245050	4	2	
20	49525029YH	40	M	1.124884	3	2	
21	25740R14MI	63	M	-0.930973	4	4	
22	84250L43IB	47	F	0.718028	3	3	
23	84258G83EV	28	F	-0.903102	2	1	
24	53254B39QD	55	F	1.799328	1	4	
25	97021A36QC	40	F	-0.358081	1	2	
26	47560Z98KV	44	F	0.465200	3	2	
27	73529T58NE	64	M	-2.017858	3	4	
28	30806A42VM	44	F	0.198266	1	2	
29	00049S85RE	43	F	0.555001	1	2	
30	26721D43SM	39	M	0.581926	4	2	

31	82508W18LT	29	F	-0.019113	1	1
32	80438M73AV	63	M	-0.719885	1	4
33	67814A83LG	51	M	-0.175610	2	3
34	87070V97Z0	33	F	-0.657252	3	1
35	79093X41WL	45	F	-0.669774	1	3
36	80170A41DA	58	M	1.649910	2	4
37	92814K57MF	35	M	-0.345982	4	2
38	67811V23ZL	63	F	-0.589992	4	4
39	49928Y380E	62	F	-0.580599	4	4
40	60766Z460D	61	F	0.699457	2	4
41	65221A93PY	54	M	-0.735907	2	3
42	67305F19B0	47	M	-0.358554	3	3
43	11497L25KD	56	F	-0.674976	1	4
44	06436U96WN	33	F	-1.114174	3	1
45	00974072DZ	59	F	-1.016883	1	4
46	09160K49TY	52	F	1.932804	1	3
47	75377M96JT	43	F	0.645703	3	2
48	50391052JL	27	F	-0.596844	1	1
49	75750M01TD	40	M	0.165158	3	2

	gender_numeric	age_car	age_car_segment	income_standardized
0	0	443	3	0.692730
1	0	335	2	1.327512
2	0	706	4	-0.543383
3	1	706	4	-0.106042
4	1	161	1	-1.001910
5	1	225	2	-0.242212
6	1	288	2	-0.180940
7	1	429	3	0.834322
8	0	521	4	1.444305
9	1	440	3	-2.093765
10	0	62	1	-1.507332
11	1	577	4	0.137720
12	1	15	1	-1.061434
13	1	65	1	-0.434468
14	1	84	1	0.685243
15	1	400	3	1.024824
16	0	37	1	-1.199001
17	1	188	1	-0.936500
18	1	65	1	-0.389229
19	1	989	4	-0.245050
20	0	476	3	1.124884
21	0	712	4	-0.930973
22	1	429	3	0.718028
23	1	260	2	-0.903102
24	1	64	1	1.799328
25	1	13	1	-0.358081
26	1	445	3	0.465200
27	0	409	3	-2.017858

28	1	153	1	0.198266
29	1	-121	1	0.555001
30	0	861	4	0.581926
31	1	-61	1	-0.019113
32	0	84	1	-0.719885
33	0	304	2	-0.175610
34	1	409	3	-0.657252
35	1	41	1	-0.669774
36	0	301	2	1.649910
37	0	783	4	-0.345982
38	1	623	4	-0.589992
39	1	643	4	-0.580599
40	1	211	2	0.699457
41	0	350	2	-0.735907
42	0	446	3	-0.358554
43	1	102	1	-0.674976
44	1	423	3	-1.114174
45	1	131	1	-1.016883
46	1	98	1	1.932804
47	1	428	3	0.645703
48	1	92	1	-0.596844
49	0	477	3	0.165158

```
# Count the number of rows in the indian_df
```

```
total_count = indian_df.shape[0]
```

```
# Display the count
```

```
print(f'Total number of records in indian_df: {total_count}')
```

```
# Optionally, you can also count the number of records for each column
```

```
data_count = indian_df.count()
```

```
# Display the count for each column
```

```
print("\nCount of non-null entries in each column:")
```

```
print(data_count)
```

```
Total number of records in indian_df: 70000
```

```
Count of non-null entries in each column:
```

```
ID          70000
CURR_AGE     70000
GENDER       70000
ANN_INCOME   70000
DT_MAINT     70000
age_group    70000
gender_numeric 70000
age_car      70000
age_car_segment 70000
income_standardized 70000
dtype: int64
```

```

# Defined features and target variables for the Japanese dataset
X_japanese = japanese_df[['CURR_AGE', 'gender_numeric',
'ANN_INCOME_INR', 'age_car_segment']]
y_japanese = japanese_df['PURCHASE']

# Split the Japanese dataset
X_train_japanese, X_test_japanese, y_train_japanese, y_test_japanese =
train_test_split(X_japanese, y_japanese, test_size=0.3,
random_state=42)

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit and transform the training data, and transform the testing data
X_train_japanese = scaler.fit_transform(X_train_japanese)
X_test_japanese = scaler.transform(X_test_japanese)

# Initialize the Logistic Regression model
model_japanese = LogisticRegression()

# Train the model
model_japanese.fit(X_train_japanese, y_train_japanese)

# Predict on the test set
y_pred_japanese = model_japanese.predict(X_test_japanese)

# Evaluate the model
print("Japanese Dataset Classification Report:")
print(classification_report(y_test_japanese, y_pred_japanese))
print("Accuracy:", accuracy_score(y_test_japanese, y_pred_japanese))
print("Confusion Matrix:")
print(confusion_matrix(y_test_japanese, y_pred_japanese))
print("ROC AUC Score:", roc_auc_score(y_test_japanese,
model_japanese.predict_proba(X_test_japanese)[: , 1]))

```

Japanese Dataset Classification Report:

	precision	recall	f1-score	support
0	0.64	0.60	0.62	5013
1	0.72	0.76	0.74	6987
accuracy			0.69	12000
macro avg	0.68	0.68	0.68	12000
weighted avg	0.69	0.69	0.69	12000

Accuracy: 0.69

Confusion Matrix:

[[2984 2029]

```

[1691 5296]]
ROC AUC Score: 0.7500105422195408

# Get feature importance
importance = model_japanese.coef_[0]
feature_names = X_japanese.columns

# Create a DataFrame for better visualization
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
importance})
importance_df = importance_df.sort_values(by='Importance',
ascending=False)
print(importance_df)

```

	Feature	Importance
3	age_car_segment	0.872753
2	ANN_INCOME_INR	0.423752
0	CURR_AGE	-0.117223
1	gender_numeric	-0.118834

```

# Import necessary libraries
from sklearn.metrics import precision_score, recall_score,
accuracy_score, roc_auc_score

# Evaluate model performance at different cutoffs
cutoffs = np.arange(0.01, 1, 0.01) # Define cutoff values from 0.01
to 0.99
accuracy_jp = []
precision_jp = []
recall_jp = []
roc_auc_jp = []

# Loop through each cutoff value and calculate performance metrics
for cutoff in cutoffs:
    y_pred_jp = np.where(y_probs_jp > cutoff, 1, 0) # Assign class
labels based on cutoff

    acc = accuracy_score(y_test_jp, y_pred_jp)
    prec = precision_score(y_test_jp, y_pred_jp, zero_division=0) #
Avoid undefined precision
    rec = recall_score(y_test_jp, y_pred_jp)
    roc_auc = roc_auc_score(y_test_jp, y_probs_jp)

    accuracy_jp.append(acc)
    precision_jp.append(prec)
    recall_jp.append(rec)
    roc_auc_jp.append(roc_auc)

# Print the cutoff value and corresponding performance metrics
print(f"Cutoff: {cutoff:.2f} | Accuracy: {acc:.4f} | Precision:

```

```

{prec:.4f} | Recall: {rec:.4f} | ROC-AUC: {roc_auc:.4f}")

# Find the best cutoff based on desired metric (e.g., accuracy)
best_cutoff_index_jp = np.argmax(accuracy_jp) # You can change this
to precision/recall depending on the metric
best_cutoff_jp = cutoffs[best_cutoff_index_jp]

print(f"\nBest cutoff value based on accuracy for Japanese dataset:
{best_cutoff_jp:.2f}")

Cutoff: 0.01 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.02 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.03 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.04 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.05 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.06 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.07 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.08 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.09 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.10 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.11 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.12 | Accuracy: 0.5823 | Precision: 0.5823 | Recall: 1.0000 |
ROC-AUC: 0.7493
Cutoff: 0.13 | Accuracy: 0.5824 | Precision: 0.5824 | Recall: 0.9997 |
ROC-AUC: 0.7493
Cutoff: 0.14 | Accuracy: 0.5847 | Precision: 0.5837 | Recall: 0.9994 |
ROC-AUC: 0.7493
Cutoff: 0.15 | Accuracy: 0.5877 | Precision: 0.5855 | Recall: 0.9989 |
ROC-AUC: 0.7493
Cutoff: 0.16 | Accuracy: 0.5906 | Precision: 0.5873 | Recall: 0.9981 |
ROC-AUC: 0.7493
Cutoff: 0.17 | Accuracy: 0.5934 | Precision: 0.5891 | Recall: 0.9970 |
ROC-AUC: 0.7493
Cutoff: 0.18 | Accuracy: 0.5958 | Precision: 0.5907 | Recall: 0.9963 |
ROC-AUC: 0.7493
Cutoff: 0.19 | Accuracy: 0.5979 | Precision: 0.5922 | Recall: 0.9934 |
ROC-AUC: 0.7493
Cutoff: 0.20 | Accuracy: 0.6012 | Precision: 0.5944 | Recall: 0.9923 |
ROC-AUC: 0.7493

```


Cutoff: 0.21		Accuracy: 0.6045		Precision: 0.5966		Recall: 0.9901	
ROC-AUC: 0.7493							
Cutoff: 0.22		Accuracy: 0.6102		Precision: 0.6005		Recall: 0.9874	
ROC-AUC: 0.7493							
Cutoff: 0.23		Accuracy: 0.6152		Precision: 0.6040		Recall: 0.9847	
ROC-AUC: 0.7493							
Cutoff: 0.24		Accuracy: 0.6174		Precision: 0.6062		Recall: 0.9785	
ROC-AUC: 0.7493							
Cutoff: 0.25		Accuracy: 0.6204		Precision: 0.6089		Recall: 0.9729	
ROC-AUC: 0.7493							
Cutoff: 0.26		Accuracy: 0.6224		Precision: 0.6110		Recall: 0.9672	
ROC-AUC: 0.7493							
Cutoff: 0.27		Accuracy: 0.6248		Precision: 0.6134		Recall: 0.9615	
ROC-AUC: 0.7493							
Cutoff: 0.28		Accuracy: 0.6255		Precision: 0.6150		Recall: 0.9541	
ROC-AUC: 0.7493							
Cutoff: 0.29		Accuracy: 0.6306		Precision: 0.6193		Recall: 0.9485	
ROC-AUC: 0.7493							
Cutoff: 0.30		Accuracy: 0.6338		Precision: 0.6226		Recall: 0.9420	
ROC-AUC: 0.7493							
Cutoff: 0.31		Accuracy: 0.6381		Precision: 0.6261		Recall: 0.9393	
ROC-AUC: 0.7493							
Cutoff: 0.32		Accuracy: 0.6401		Precision: 0.6284		Recall: 0.9346	
ROC-AUC: 0.7493							
Cutoff: 0.33		Accuracy: 0.6435		Precision: 0.6316		Recall: 0.9302	
ROC-AUC: 0.7493							
Cutoff: 0.34		Accuracy: 0.6456		Precision: 0.6339		Recall: 0.9260	
ROC-AUC: 0.7493							
Cutoff: 0.35		Accuracy: 0.6476		Precision: 0.6361		Recall: 0.9226	
ROC-AUC: 0.7493							
Cutoff: 0.36		Accuracy: 0.6512		Precision: 0.6396		Recall: 0.9181	
ROC-AUC: 0.7493							
Cutoff: 0.37		Accuracy: 0.6538		Precision: 0.6427		Recall: 0.9127	
ROC-AUC: 0.7493							
Cutoff: 0.38		Accuracy: 0.6584		Precision: 0.6475		Recall: 0.9071	
ROC-AUC: 0.7493							
Cutoff: 0.39		Accuracy: 0.6628		Precision: 0.6524		Recall: 0.9011	
ROC-AUC: 0.7493							
Cutoff: 0.40		Accuracy: 0.6675		Precision: 0.6578		Recall: 0.8941	
ROC-AUC: 0.7493							
Cutoff: 0.41		Accuracy: 0.6719		Precision: 0.6637		Recall: 0.8849	
ROC-AUC: 0.7493							
Cutoff: 0.42		Accuracy: 0.6750		Precision: 0.6694		Recall: 0.8728	
ROC-AUC: 0.7493							
Cutoff: 0.43		Accuracy: 0.6793		Precision: 0.6764		Recall: 0.8612	
ROC-AUC: 0.7493							
Cutoff: 0.44		Accuracy: 0.6827		Precision: 0.6831		Recall: 0.8490	
ROC-AUC: 0.7493							
Cutoff: 0.45		Accuracy: 0.6845		Precision: 0.6893		Recall: 0.8343	

ROC-AUC: 0.7493			
Cutoff: 0.46	Accuracy: 0.6831	Precision: 0.6934	Recall: 0.8169
ROC-AUC: 0.7493			
Cutoff: 0.47	Accuracy: 0.6844	Precision: 0.6993	Recall: 0.8035
ROC-AUC: 0.7493			
Cutoff: 0.48	Accuracy: 0.6855	Precision: 0.7055	Recall: 0.7895
ROC-AUC: 0.7493			
Cutoff: 0.49	Accuracy: 0.6884	Precision: 0.7145	Recall: 0.7742
ROC-AUC: 0.7493			
Cutoff: 0.50	Accuracy: 0.6913	Precision: 0.7237	Recall: 0.7598
ROC-AUC: 0.7493			
Cutoff: 0.51	Accuracy: 0.6903	Precision: 0.7310	Recall: 0.7407
ROC-AUC: 0.7493			
Cutoff: 0.52	Accuracy: 0.6880	Precision: 0.7377	Recall: 0.7202
ROC-AUC: 0.7493			
Cutoff: 0.53	Accuracy: 0.6888	Precision: 0.7453	Recall: 0.7073
ROC-AUC: 0.7493			
Cutoff: 0.54	Accuracy: 0.6898	Precision: 0.7531	Recall: 0.6951
ROC-AUC: 0.7493			
Cutoff: 0.55	Accuracy: 0.6897	Precision: 0.7588	Recall: 0.6846
ROC-AUC: 0.7493			
Cutoff: 0.56	Accuracy: 0.6880	Precision: 0.7646	Recall: 0.6705
ROC-AUC: 0.7493			
Cutoff: 0.57	Accuracy: 0.6881	Precision: 0.7705	Recall: 0.6612
ROC-AUC: 0.7493			
Cutoff: 0.58	Accuracy: 0.6866	Precision: 0.7745	Recall: 0.6514
ROC-AUC: 0.7493			
Cutoff: 0.59	Accuracy: 0.6853	Precision: 0.7791	Recall: 0.6415
ROC-AUC: 0.7493			
Cutoff: 0.60	Accuracy: 0.6836	Precision: 0.7827	Recall: 0.6320
ROC-AUC: 0.7493			
Cutoff: 0.61	Accuracy: 0.6817	Precision: 0.7878	Recall: 0.6206
ROC-AUC: 0.7493			
Cutoff: 0.62	Accuracy: 0.6793	Precision: 0.7925	Recall: 0.6084
ROC-AUC: 0.7493			
Cutoff: 0.63	Accuracy: 0.6770	Precision: 0.7980	Recall: 0.5961
ROC-AUC: 0.7493			
Cutoff: 0.64	Accuracy: 0.6736	Precision: 0.8025	Recall: 0.5828
ROC-AUC: 0.7493			
Cutoff: 0.65	Accuracy: 0.6676	Precision: 0.8070	Recall: 0.5639
ROC-AUC: 0.7493			
Cutoff: 0.66	Accuracy: 0.6606	Precision: 0.8112	Recall: 0.5436
ROC-AUC: 0.7493			
Cutoff: 0.67	Accuracy: 0.6521	Precision: 0.8138	Recall: 0.5218
ROC-AUC: 0.7493			
Cutoff: 0.68	Accuracy: 0.6458	Precision: 0.8201	Recall: 0.5018
ROC-AUC: 0.7493			
Cutoff: 0.69	Accuracy: 0.6370	Precision: 0.8236	Recall: 0.4792
ROC-AUC: 0.7493			

Cutoff: 0.70	Accuracy: 0.6285	Precision: 0.8301	Recall: 0.4551
ROC-AUC: 0.7493			
Cutoff: 0.71	Accuracy: 0.6184	Precision: 0.8337	Recall: 0.4305
ROC-AUC: 0.7493			
Cutoff: 0.72	Accuracy: 0.6098	Precision: 0.8398	Recall: 0.4075
ROC-AUC: 0.7493			
Cutoff: 0.73	Accuracy: 0.5991	Precision: 0.8474	Recall: 0.3798
ROC-AUC: 0.7493			
Cutoff: 0.74	Accuracy: 0.5873	Precision: 0.8536	Recall: 0.3514
ROC-AUC: 0.7493			
Cutoff: 0.75	Accuracy: 0.5777	Precision: 0.8582	Recall: 0.3290
ROC-AUC: 0.7493			
Cutoff: 0.76	Accuracy: 0.5701	Precision: 0.8662	Recall: 0.3094
ROC-AUC: 0.7493			
Cutoff: 0.77	Accuracy: 0.5622	Precision: 0.8717	Recall: 0.2908
ROC-AUC: 0.7493			
Cutoff: 0.78	Accuracy: 0.5541	Precision: 0.8794	Recall: 0.2714
ROC-AUC: 0.7493			
Cutoff: 0.79	Accuracy: 0.5463	Precision: 0.8836	Recall: 0.2542
ROC-AUC: 0.7493			
Cutoff: 0.80	Accuracy: 0.5377	Precision: 0.8842	Recall: 0.2372
ROC-AUC: 0.7493			
Cutoff: 0.81	Accuracy: 0.5293	Precision: 0.8828	Recall: 0.2210
ROC-AUC: 0.7493			
Cutoff: 0.82	Accuracy: 0.5228	Precision: 0.8814	Recall: 0.2084
ROC-AUC: 0.7493			
Cutoff: 0.83	Accuracy: 0.5155	Precision: 0.8876	Recall: 0.1922
ROC-AUC: 0.7493			
Cutoff: 0.84	Accuracy: 0.5060	Precision: 0.8972	Recall: 0.1712
ROC-AUC: 0.7493			
Cutoff: 0.85	Accuracy: 0.4944	Precision: 0.9000	Recall: 0.1481
ROC-AUC: 0.7493			
Cutoff: 0.86	Accuracy: 0.4823	Precision: 0.9006	Recall: 0.1245
ROC-AUC: 0.7493			
Cutoff: 0.87	Accuracy: 0.4707	Precision: 0.9034	Recall: 0.1018
ROC-AUC: 0.7493			
Cutoff: 0.88	Accuracy: 0.4599	Precision: 0.9161	Recall: 0.0797
ROC-AUC: 0.7493			
Cutoff: 0.89	Accuracy: 0.4528	Precision: 0.9218	Recall: 0.0658
ROC-AUC: 0.7493			
Cutoff: 0.90	Accuracy: 0.4478	Precision: 0.9188	Recall: 0.0567
ROC-AUC: 0.7493			
Cutoff: 0.91	Accuracy: 0.4449	Precision: 0.9179	Recall: 0.0512
ROC-AUC: 0.7493			
Cutoff: 0.92	Accuracy: 0.4397	Precision: 0.9122	Recall: 0.0416
ROC-AUC: 0.7493			
Cutoff: 0.93	Accuracy: 0.4343	Precision: 0.9160	Recall: 0.0312
ROC-AUC: 0.7493			
Cutoff: 0.94	Accuracy: 0.4278	Precision: 0.9054	Recall: 0.0192

```
ROC-AUC: 0.7493
Cutoff: 0.95 | Accuracy: 0.4206 | Precision: 0.9250 | Recall: 0.0053 |
ROC-AUC: 0.7493
Cutoff: 0.96 | Accuracy: 0.4178 | Precision: 0.0000 | Recall: 0.0000 |
ROC-AUC: 0.7493
Cutoff: 0.97 | Accuracy: 0.4178 | Precision: 0.0000 | Recall: 0.0000 |
ROC-AUC: 0.7493
Cutoff: 0.98 | Accuracy: 0.4178 | Precision: 0.0000 | Recall: 0.0000 |
ROC-AUC: 0.7493
Cutoff: 0.99 | Accuracy: 0.4178 | Precision: 0.0000 | Recall: 0.0000 |
ROC-AUC: 0.7493
```

Best cutoff value based on accuracy for Japanese dataset: 0.50

```
# Prepare your features and target variable
X_jp = japanese_df[['age_group', 'gender_numeric', 'age_car_segment',
'income_standardized']]
y_jp = japanese_df['PURCHASE'] # Replace with your target column name

# Split the dataset into training and testing sets
X_train_jp, X_test_jp, y_train_jp, y_test_jp = train_test_split(X_jp,
y_jp, test_size=0.2, random_state=42)

# Create and fit the logistic regression model
model_jp = LogisticRegression()
model_jp.fit(X_train_jp, y_train_jp)

# Get the intercept and coefficients
intercept_jp = model_jp.intercept_[0] # Get the intercept
coefficients_jp = model_jp.coef_[0] # Get the coefficients for each
feature

# Create a dictionary to display the intercept and coefficients
coefficients_dict_jp = {
    'Intercept': intercept_jp,
    'age_group': coefficients_jp[0],
    'gender_numeric': coefficients_jp[1],
    'age_car_segment': coefficients_jp[2],
    'income_standardized': coefficients_jp[3]
}

# Display the results
coefficients_df_jp = pd.DataFrame(coefficients_dict_jp, index=[0])
print("Intercept and Coefficients for the Japanese Dataset:")
print(coefficients_df_jp)

Intercept and Coefficients for the Japanese Dataset:
Intercept age_group gender_numeric age_car_segment
income_standardized
```

```
0    -1.573318    -0.099991        -0.229322        0.96041
0.419205
```

```
import numpy as np

# Coefficients from the logistic regression model trained on the
# Japanese dataset
intercept_jp = -1.573318
coeff_age_group_jp = -0.099991
coeff_gender_jp = -0.229322
coeff_age_car_segment_jp = 0.96041
coeff_income_standardized_jp = 0.419205

# Function to calculate probability for each row using coefficients
# from the Japanese dataset
def calculate_purchase_probability_jp(age_group, gender_numeric,
income_standardized, age_car_segment):
    # Logistic regression equation
    logit = (intercept_jp
              + coeff_age_group_jp * age_group
              + coeff_gender_jp * gender_numeric
              + coeff_income_standardized_jp * income_standardized
              + coeff_age_car_segment_jp * age_car_segment)

    # Apply logistic function
    probability = 1 / (1 + np.exp(-logit))

    return round(probability, 2) # Round to 2 decimal places for
better precision

# Apply the function to the entire Indian dataset (indian_df)
indian_df['purchase_probability'] = indian_df.apply(
    lambda row: calculate_purchase_probability_jp(row['age_group'],
row['gender_numeric'],
row['income_standardized'],
row['age_car_segment']),
    axis=1
)

# View the first 5 rows of the updated dataframe
print(indian_df.head(5))

# Optionally, save the updated dataframe with probabilities to Excel
indian_df.to_excel("indian_df_with_purchase_probabilities.xlsx",
index=False)
```

	ID	CURR_AGE	GENDER	ANN_INCOME	DT_MAINT	age_group	\
0	20710B05XL	54	M	0.692730	3	3	
1	89602T51HX	47	M	1.327512	2	3	
2	70190Z52IP	60	M	-0.543383	4	4	
3	25623V15MU	55	F	-0.106042	4	4	
4	36230I68CE	32	F	-1.001910	1	1	

	gender_numeric	age_car	age_car_segment	income_standardized	\
0	0	443	3	0.692730	
1	0	335	2	1.327512	
2	0	706	4	-0.543383	
3	1	706	4	-0.106042	
4	1	161	1	-1.001910	

	purchase_probability	purchase
0	0.79	1
1	0.65	1
2	0.84	1
3	0.83	1
4	0.20	1

Apply the if-else condition to assign 1 or 0 based on the cutoff value of 0.2

```
indian_df['purchase'] = indian_df['purchase_probability'].apply(lambda x: 1 if x > 0.2 else 0)
```

Show the first 50 rows with the new 'purchase' column

```
print(indian_df[['CURR_AGE', 'gender_numeric', 'ANN_INCOME', 'age_car_segment', 'purchase_probability', 'purchase']].head(50))
```

Optionally, save the updated dataframe to Excel

```
indian_df.to_excel("indian_df_with_purchase_predictions.xlsx", index=False)
```

	CURR_AGE	gender_numeric	ANN_INCOME	age_car_segment	\
0	54	0	0.692730	3	
1	47	0	1.327512	2	
2	60	0	-0.543383	4	
3	55	1	-0.106042	4	
4	32	1	-1.001910	1	
5	48	1	-0.242212	2	
6	26	1	-0.180940	2	
7	45	1	0.834322	3	
8	55	0	1.444305	4	
9	64	1	-2.093765	3	
10	53	0	-1.507332	1	
11	44	1	0.137720	4	
12	59	1	-1.061434	1	
13	27	1	-0.434468	1	

14	57	1	0.685243	1
15	40	1	1.024824	3
16	33	0	-1.199001	1
17	57	1	-0.936500	1
18	59	1	-0.389229	1
19	42	1	-0.245050	4
20	40	0	1.124884	3
21	63	0	-0.930973	4
22	47	1	0.718028	3
23	28	1	-0.903102	2
24	55	1	1.799328	1
25	40	1	-0.358081	1
26	44	1	0.465200	3
27	64	0	-2.017858	3
28	44	1	0.198266	1
29	43	1	0.555001	1
30	39	0	0.581926	4
31	29	1	-0.019113	1
32	63	0	-0.719885	1
33	51	0	-0.175610	2
34	33	1	-0.657252	3
35	45	1	-0.669774	1
36	58	0	1.649910	2
37	35	0	-0.345982	4
38	63	1	-0.589992	4
39	62	1	-0.580599	4
40	61	1	0.699457	2
41	54	0	-0.735907	2
42	47	0	-0.358554	3
43	56	1	-0.674976	1
44	33	1	-1.114174	3
45	59	1	-1.016883	1
46	52	1	1.932804	1
47	43	1	0.645703	3
48	27	1	-0.596844	1
49	40	0	0.165158	3

	purchase_probability	purchase
0	0.79	1
1	0.65	1
2	0.84	1
3	0.83	1
4	0.20	0
5	0.43	1
6	0.49	1
7	0.76	1
8	0.92	1
9	0.45	1
10	0.18	0

11	0.87	1
12	0.16	0
13	0.25	1
14	0.28	1
15	0.79	1
16	0.23	1
17	0.16	0
18	0.20	0
19	0.85	1
20	0.83	1
21	0.81	1
22	0.75	1
23	0.41	1
24	0.38	1
25	0.23	1
26	0.75	1
27	0.52	1
28	0.28	1
29	0.31	1
30	0.91	1
31	0.28	1
32	0.21	1
33	0.49	1
34	0.67	1
35	0.19	0
36	0.65	1
37	0.87	1
38	0.80	1
39	0.80	1
40	0.50	1
41	0.44	1
42	0.70	1
43	0.18	0
44	0.63	1
45	0.16	0
46	0.42	1
47	0.76	1
48	0.23	1
49	0.76	1

```
# Count the occurrences of 0 and 1 in the 'purchase' column
purchase_counts = indian_df['purchase'].value_counts()
```

```
# Print the count of 0 and 1
```

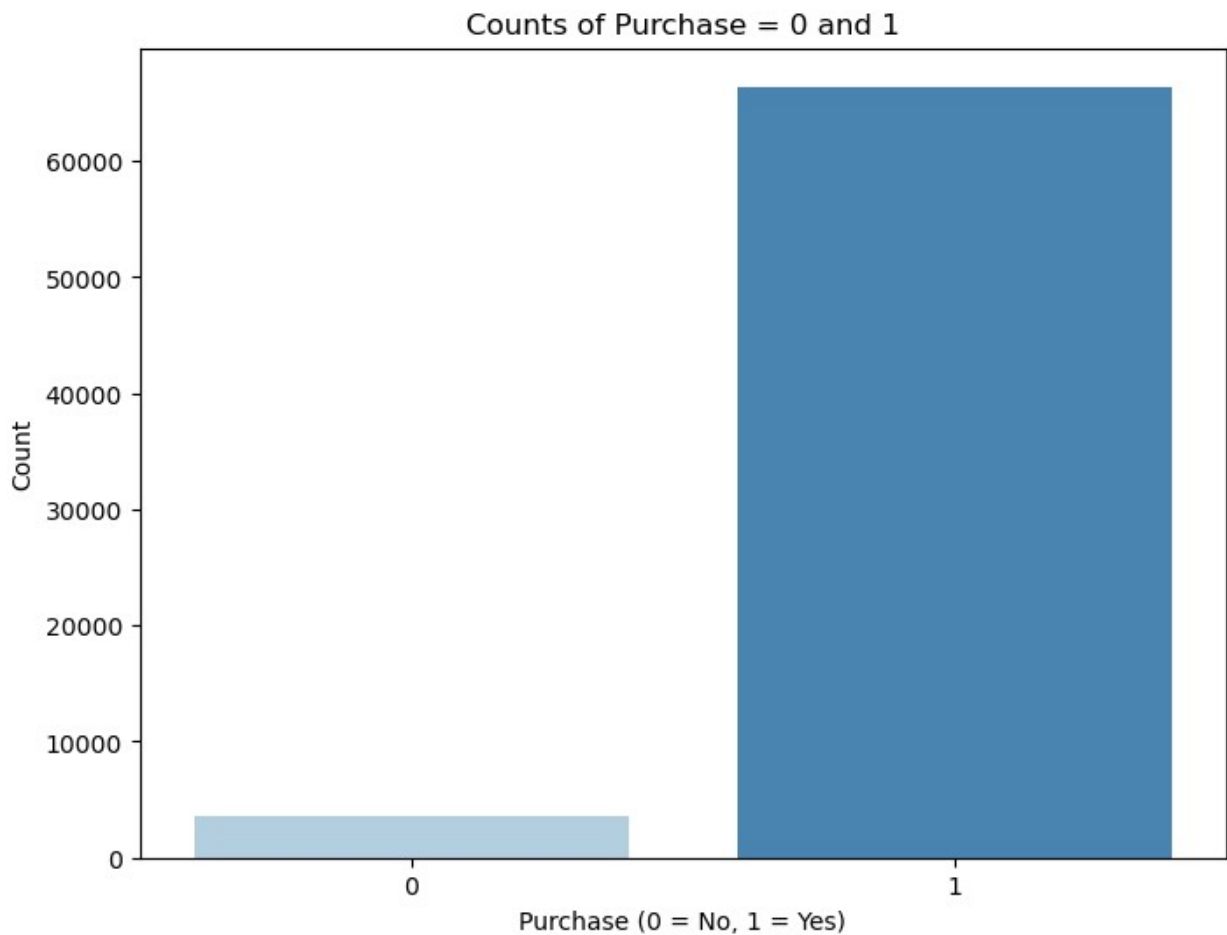
```
print("Count of Purchase = 0:", purchase_counts[0])
print("Count of Purchase = 1:", purchase_counts[1])
```

```
Count of Purchase = 0: 3601
Count of Purchase = 1: 66399
```




```
import matplotlib.pyplot as plt
import seaborn as sns

# Bar plot of purchase counts
plt.figure(figsize=(8, 6))
sns.countplot(x='purchase', data=indian_df, palette='Blues')
plt.title('Counts of Purchase = 0 and 1')
plt.xlabel('Purchase (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.show()
```



```
# Pie chart for purchase proportions
purchase_counts = indian_df['purchase'].value_counts()
plt.figure(figsize=(8, 8))
plt.pie(purchase_counts, labels=['No Purchase (0)', 'Purchase (1)'],
autopct='%1.1f%%', startangle=90, colors=['skyblue', 'lightgreen'])
plt.title('Proportion of Purchase Predictions')
plt.show()
```

Proportion of Purchase Predictions



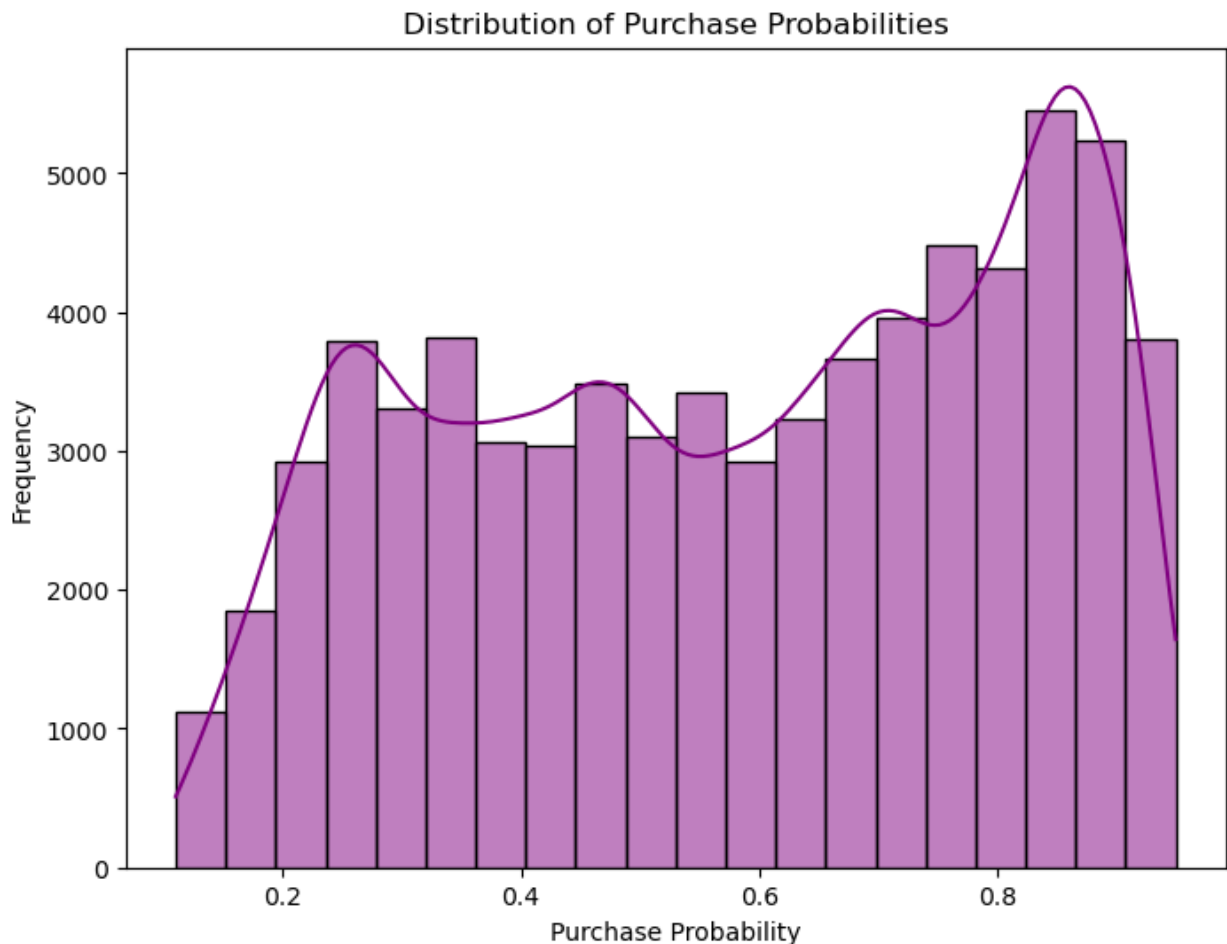
```
# Histogram of purchase probabilities
plt.figure(figsize=(8, 6))
sns.histplot(indian_df['purchase_probability'], bins=20, kde=True,
color='purple')
plt.title('Distribution of Purchase Probabilities')
plt.xlabel('Purchase Probability')
plt.ylabel('Frequency')
plt.show()

# Replace infinite values in the dataset with NaN
indian_df.replace([np.inf, -np.inf], np.nan, inplace=True)
```

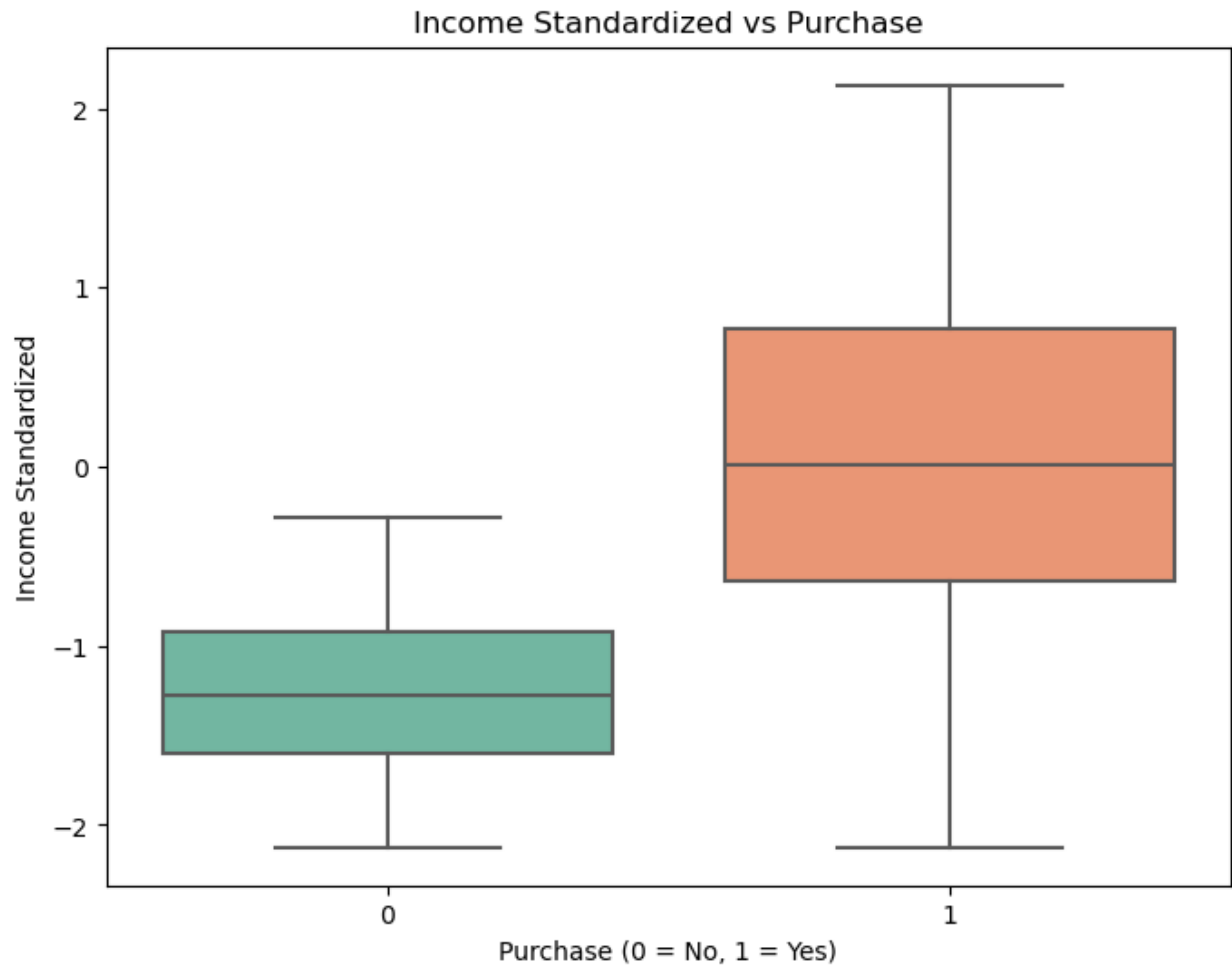
```
# After replacing, you might also want to drop rows with NaN values or
fill them:
# Drop rows with NaN values (optional)
indian_df.dropna(inplace=True)
```

```
C:\Users\chakr\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```



```
# Box plot of income_standardized vs purchase
plt.figure(figsize=(8, 6))
sns.boxplot(x='purchase', y='income_standardized', data=indian_df,
palette='Set2')
plt.title('Income Standardized vs Purchase')
plt.xlabel('Purchase (0 = No, 1 = Yes)')
plt.ylabel('Income Standardized')
plt.show()
```



```
# Box plot of age_car_segment vs purchase
plt.figure(figsize=(8, 6))
sns.boxplot(x='purchase', y='age_car_segment', data=indian_df,
palette='Set1')
plt.title('Age Car Segment vs Purchase')
plt.xlabel('Purchase (0 = No, 1 = Yes)')
plt.ylabel('Age Car Segment')
plt.show()
```

